

Mixtral-8x7B: Parallelism Comparison on
4 A100 GPUs for Batch Size = 64

